

Adrij Shikhar

Personal Information

Email
adrijshikhar85@gmail.com

Phone
+918218058928

Github
<https://github.com/adrijshikhar>

LinkedIn
<https://www.linkedin.com/in/adrij-shikhar/>

Website
<https://adrijshikhar.dev>

Education

July 2018 - May 2022
Indian Institute of Technology, Roorkee
Chemical B.Tech.

Work History

Senior Software Engineer

Jun 2022 - Current

Hevo Data, Bengaluru

Agent-Native Connector Development

- *Connector Agent*: Built the surface that lets LLMs generate and operate connectors — a construct library, a connector generator, an installable CLI, a skill library, and an MCP server exposing platform APIs to agents.
- *Tooling Platform*: Started the internal agent-tooling platform the team now builds on — repo structure, skill conventions, CI auto-release, marketplace distribution.

Connector Platform & Developer Experience

- *Declarative Auth*: Made connector authentication declarative. A provider used to be its own processor inside the platform — a code change, a database migration, and coordination across three services. It is now a config file, with no platform change at all.
- *Ownership Boundary*: Moved runtime configuration and config templates into the connector, behind a compatibility kit that gates correctness at build time instead of in review.
- *Fleet Standardisation*: Standardised how connector repos are created, versioned and released behind a central BOM, and decoupled the fleet from the shared platform SDK.

Connection Reliability

- *Test-Connection Contract*: Owned the contract end to end and took its false-negative rate from over 30% to nil in validation — a failing check now means an actual connection problem.
- *Validation Coverage*: Delivered configuration validation across every source and destination — Postgres, MySQL, Oracle, SQL Server, Snowflake, BigQuery, Redshift — so a misconfiguration surfaces at connect time rather than at first sync, with p95 latency under 10 seconds.

Failure Visibility & Diagnosis

- *Dynamic Classification*: Replaced static error classification with a dynamic system. A classification change used to wait on a connector's full release cycle; it now takes minutes instead of 3–7 days.
- *Regex Engine Root Cause*: Root-caused a class of production hangs to the regex engine rather than the queries being run, and migrated to RE2J — bounding worst-case evaluation under five seconds and taking a load run from a multi-minute stall to 1m19s.
- *Observability*: Built metrics and dashboards for connection health, and put HTTP client metrics into the connector construct library so teams see a rate-limiting API before it becomes a sync failure.

Data Correctness at Scale

- *History Mode*: Delivered SCD Type 2 history mode across Snowflake, BigQuery and Redshift, handling the per-destination SQL differences — Redshift has no MERGE and needs a four-statement path.
- *CDC Framework*: Implemented real-time CDC on the Debezium engine and built the schema catalog service behind it from scratch — schema versioning, metadata management and compatibility checks, so a pipeline survives schema evolution.
- *Object Scale*: Scaled object handling from 1,000 to 25,000 objects with a 5x gain in MySQL ingestion throughput, and added parent-child handling with inferred deletes.

Engineering Velocity

- *Build & CI Runtime*: Cut runtime with measured before and after — `mvn clean verify` 31:43 to 10:33 and CI 43:32 to 21:24 on the first service, then 11:07 to 5:15 as the approach was adopted elsewhere. Sweeps run against a 90-day baseline, not by feel.
- *Local Development*: Built the unified CLI and local development platform so the team runs the stack locally instead of queueing for shared environments.

References

Tariq Iqbal | Senior Engineering Manager - Atlassian
[linkedin.com/in/ta-riq](https://www.linkedin.com/in/ta-riq)

Shubham Goyal | CTO - Powerplay
[linkedin.com/in/shubham7stark](https://www.linkedin.com/in/shubham7stark)